

Q775 Assignment 1

Name: Ahana Malhotra

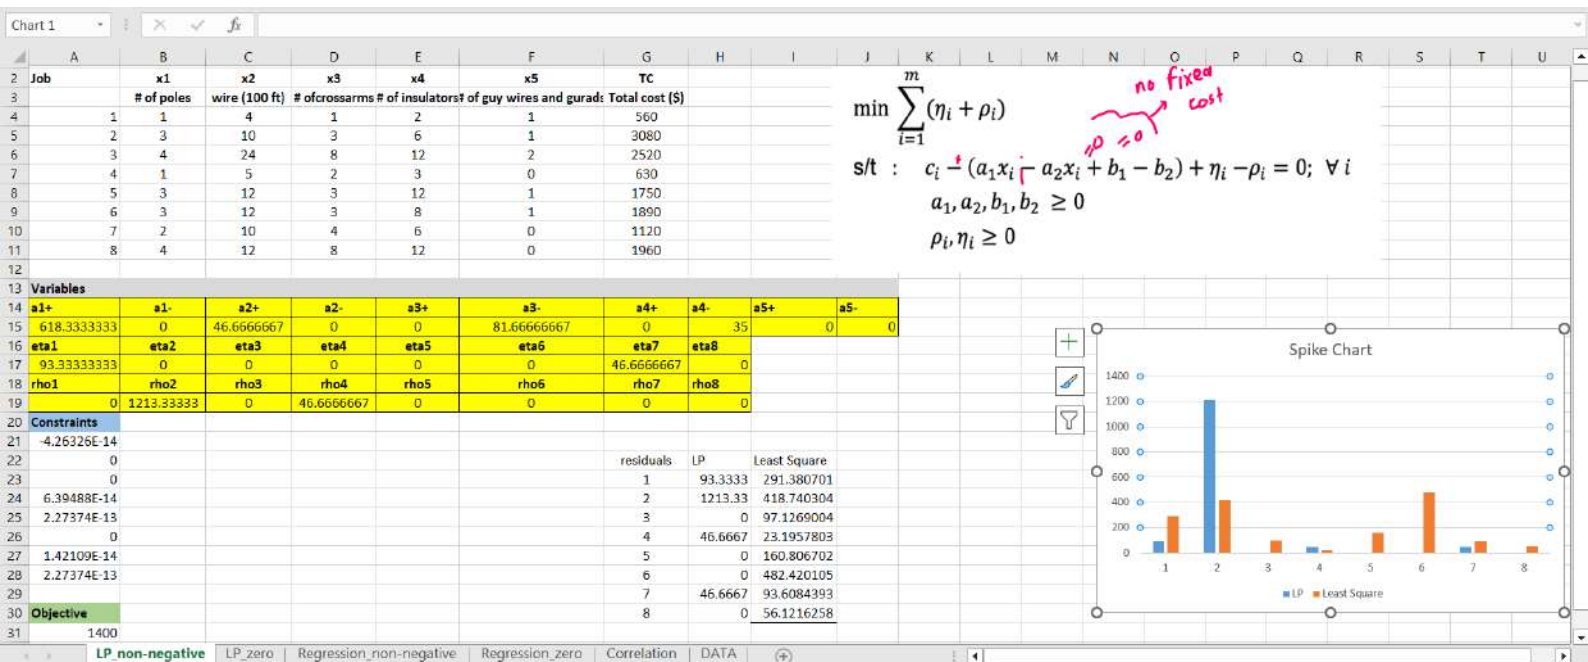
Student ID:400380436

Date of Submission: 10 March 2023


1. **Develop the CER prediction model used in class to force all the coefficients of the function to be non-negative. Compare the resulting residuals with those found in class and comment on the results.**

CER prediction model is developed and two cases are performed. In the first case, all the coefficients of the functions to be non-negative are enforced and in second case all the coefficients of the functions can be non-negative. When we analyzed the residuals, we found when non-negative coefficients are not enforced, then the value of residuals have decreased, which indicates better accuracy of prediction of total cost.

Below sheet stating 'LP-non negative' means non-negative coefficients are enforced and in others they are not.



M22														
	A	B	C	D	E	F	G	H	I	J	K	L	M	N
1	SUMMARY OUTPUT													
2										RESIDUAL OUTPUT				
3	Regression Statistics													
4	Multiple R	0.99039								Observations				
5	R Square	0.98088								1	851.381	-291.38		
6	Adjusted R Square	0.62205								2	2661.26	418.74		
7	Standard Error	424.546								3	2422.87	97.1269		
8	Observations	8								4	606.804	23.1958		
9										5	1589.19	160.807		
10	ANOVA													
11		df	SS	MS	F	Significance F				6	2372.42	-482.42		
12	Regression	5	2.8E+07	5547436	30.7782	0.03177				7	1213.61	-93.608		
13	Residual	3	540718	180239						8	2016.12	-56.122		
14	Total	8	2.8E+07											
15														
16		Coefficients	Standard Error	t Stat	P-value	Lower 95%	Upper 95%	Lower 95.0%	Upper 95.0%					
17	Intercept	0	#N/A	#N/A	#N/A	#N/A	#N/A	#N/A	#N/A					
18	# of poles	1347.49	349.129	3.85959	0.03074	236.411	2458.58	236.411	2458.58					
19	wire (100	51.3869	88.6669	0.57955	0.60287	-230.79	333.565	-230.79	333.565					
20	# of cross	-205.1	164.13	-1.2496	0.30004	-727.44	317.231	-727.44	317.231					
21	# of insula	-195.81	102.728	-1.9061	0.15272	-522.73	131.121	-522.73	131.121					
22	# of guy w	-104.95	467.106	-0.2247	0.83667	-1591.5	1381.59	-1591.5	1381.59					
23														
24														
25														
26														
27														
28														
29														
30														
31														
		LP_non-negative	LP_zero	Regression_non-negative	Regression_zero	Correlation	DATA							

$$\begin{aligned} \min \quad & \sum_{i=1}^n (\eta_i + \rho_i) \\ \text{s/t : } \quad & c_i + (a_1 x_i - a_2 x_i + b_1 - b_2) + \eta_i - \rho_i = 0; \quad \forall i \\ & a_1, a_2, b_1, b_2 \geq 0 \\ & \rho_i, \eta_i \geq 0 \end{aligned}$$


	A	B	C	D	E	F	G	H	I	J	K	L	M	N
1	SUMMARY OUTPUT							RESIDUAL OUTPUT						
2														
3	Regression Statistics							Observations						
4	Multiple R	0.953574						1	709.8765	-149.876				
5	R Square	0.909304						2	2690.389	389.6108				
6	Adjusted R	0.682563						3	2470.041	49.95883				
7	Standard Error	497.7263						4	475.0936	154.9064				
8	Observations	8						5	1580.174	169.826				
9								6	2399.478	-509.478				
10	ANOVA							7	1152.524	-32.5243				
11		df	SS	MS	F	Significance F		8	2032.423	-72.4233				
12	Regression	5	4967425	993484.9	4.010329	0.211553								
13	Residual	2	495463	247731.5										
14	Total	7	5462888											
15														
16		Coefficients	Standard Error	t Stat	P-value	Lower 95%	Upper 95%	Lower 95.0%	Upper 95.0%					
17	Intercept	-202.337	473.4054	-0.42741	0.7107	-2239.24	1834.562	-2239.24	1834.562					
18	# of poles	1448.538	472.6765	3.064543	0.092019	-585.225	3482.3	-585.225	3482.3					
19	wire (100 f	59.37032	105.6155	0.562136	0.630621	-395.057	513.7972	-395.057	513.7972					
20	# of crossa	-226.74	198.9694	-1.13957	0.372555	-1082.84	629.356	-1082.84	629.356					
21	# of insula	-204.826	122.2707	-1.67519	0.235883	-730.914	321.2622	-730.914	321.2622					
22	# of guy w	-137.413	552.8655	-0.24855	0.826904	-2516.2	2241.375	-2516.2	2241.375					
23														
24														
25														
26														
27														
28														
29														

	A	B	C	D	E	F	G	H
1		# of poles	wire (100 ft)	# ofcrossarms	# of insulators	# of guy wires and gurads	Total cost (\$)	
2	# of poles	1						
3	wire (100 ft)	0.818322323	1					
4	# ofcrossarms	0.826767382	0.789436128	1				
5	# of insulators	0.912732554	0.787398955	0.77770072	1			
6	# of guy wires and gurads	0.38271893	0.639516355	0.15430335	0.310349551	1		
7	Total cost (\$)	0.823202005	0.65870618	0.510091751	0.599221419	0.500268216	1	
8								
9								
10								
11								
12								
13								
14								
15								
16								
17								
18								
19								
20								
21								
22								
23								
24								
25								
26								
27								
28								
29								

2. For this question you will use the data '[general_classification_data.csv](#)'

(a) Use the Linear Discriminant Analysis approach (e.g., from Python SKLearn package or code it based on equations discussed in class) to develop a discriminant line. Comment on the results achieved and compare them to those achieved by the linear programming method we discussed in class (and implemented in Python).

Developed discriminant line from Linear Discriminant Analysis approach cannot distinguish two classes as these classes because LDA assume that the covariance matrix is same and it tries to find the maximum linear separation between classes while minimizing their class variance. So, in this case LDA could not find linear combination of features which leads to distinguishing between two classes.

However, linear programming could separate the two classes successfully, because it was able to separate the boundaries by finding optimal linear boundary. As both the techniques have different assumptions and result varies according to the input data. But linear programming can solve wide range of optimization problems as compared to LDA.

Linear Discriminant Analysis

```
In 66 1 import pandas as pd
      2 import numpy as np
```

```
In 67 1 data = pd.read_csv("./general_classification_data.csv", index_col=0)
```

```
In 68 1 data.head()
```

Out 68 ▾

< < 5 rows > > 5 rows × 3 columns				
Object	Score1	Score2	Class	
1	0.2	0.5	1	
2	0.2	0.8	1	
3	0.3	0.4	1	
4	0.3	0.7	1	
5	0.4	0.3	1	

```
In 69 1 #import os
      2 #cwd = os.getcwd()
      3 #print(cwd)
      4
```



```
In 71 1 X = data[['Score1', 'Score2']]
```

```
In 72 1 X.head()
```

Out 72 ▾

< < 5 rows ▾ > > 5 rows × 2 columns			
Object ↕	Score1 ↕	Score2 ↕	
1	0.2	0.5	
2	0.2	0.8	
3	0.3	0.4	
4	0.3	0.7	
5	0.4	0.3	

```
In 73 1 print(type(X))
      2 print(X.shape)
```

```
<class 'pandas.core.frame.DataFrame'>
(20, 2)
```

```
In 74 1 y = data.Class
      2 y.head()
```

Out 74 ▾

|< < 5 rows ▾ > >| Length: 5, dtype: int64

Object ↕	Class ↕
1	1
2	1
3	1
4	1
5	1

```
In 75 1 print(type(y))
      2 print(y.shape)
```

```
<class 'pandas.core.series.Series'>
(20,)
```


```
In 76 1 from sklearn.discriminant_analysis import LinearDiscriminantAnalysis
      2
      3 lda = LinearDiscriminantAnalysis()
      4 lda.fit(X, y)
      5
      6 coef = lda.coef_[0]
      7 intercept = lda.intercept_
      8
```

```
In 77 1 import seaborn as sns
      2 import matplotlib.pyplot as plt
      -
```

```

3
4 sns.lmplot(x='Score1', y='Score2', data=data, hue='Class', fit_reg=False)
5
6 x_vals = np.array([0, 1])
7 y_vals = -(coef[0] * x_vals + intercept) / coef[1]
8
9 plt.plot(x_vals, y_vals, 'r--', linewidth=2)
10
11 plt.show()
12

```



(b) Use the appropriate Neural Network code (MATLAB or Python) to obtain a neural network classifier for the same data set. Compare this result to that obtained in part 2(a).

Feedforward artificial neural network code has been developed in MATLAB. It classifies the dataset into two classes with accuracy of 100%. Confusion matrix has been analyzed to check the accuracy. In the confusion matrix, we can analyze that in training data, it predicted the data with 100% accuracy and so in validation and test data.

In part 2a, accuracy of prediction of classes using linear discrimination analysis was extremely low, but linear programming method helped us getting better classification among classes because it was able to find optimal linear boundary. But if data would have been non-linear, it would be difficult for linear programming as well to separate the classes and find optimal solution. So in linear and non-linear input cases, neural network gives us best accuracy and optimal results.

In case data becomes more complex, then neural network classifier will work the best among the three methods discussed in this question.

```
%% Artificial Neural Network (ANN)
```

```
clear; clc;
```

```
%% Loading the general classification data, with the classes at the last column
```

```
A = load('general_classification_data.txt');
```

```
% C is the matrix of the feature vectors
```

```
C = A(1:end,1:end-1)';
```

```
%% Create matrix for classes (targets)
```

```
T=zeros(size(unique(A(:, end)),1),size(A,1));
```

```
for i=1:size(A,1)
```

```
    T(A(i,end),i)=1;    %create identity vectors to reflect class type
```

```
end
```

```
%% Initializing a neural network 'net'
```

```
net = feedforwardnet;
```

```
net = configure(net, C, T);    % net - NN, C-feature vector, T-target
```

```
%% Setting hidden layer
```

```
hiddenLayerSize = 10; % Setting the number of hidden layers to 10
```

```
net = patternnet(hiddenLayerSize); % Pattern recognition network
```

```
%% Split data for training, validation and testing
```

```
net.divideParam.trainRatio = 0.7; % Ratio of training data is 70%
```

```
net.divideParam.valRatio = 0.2; % Ratio of validation data is 20%
```

```
net.divideParam.testRatio = 0.1; % Ratio of testing data is 10%
```

%% Training the network

[net, tr] = train(net, C, T); % the resulting model is the output net

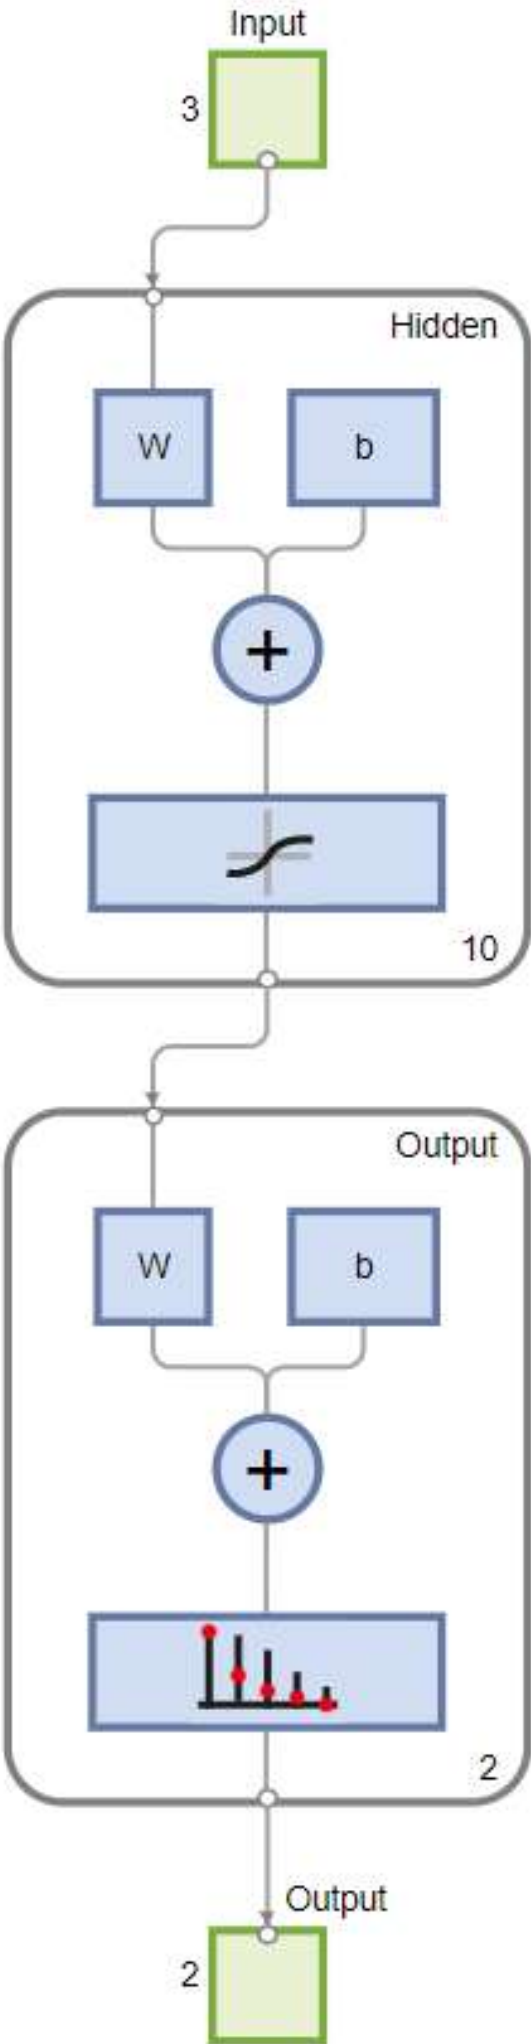
outputs = net(C);

%% Computing the classification errors

errors = gsubtract(T, outputs);

performance = perform(net, T, outputs);

view(net)



[Network Diagram](#)

Training Results

Training finished: Reached minimum gradient ✓

Training Progress

Unit	Initial Value	Stopped Value	Target Value	
Epoch	0	24	1000	▲
Elapsed Time	-	00:00:00	-	
Performance	0.206	2.74e-07	0	
Gradient	0.347	6.86e-07	1e-06	
Validation Checks	0	1	6	▼

Training Algorithms

Data Division: Random dividerand

Training: Scaled Conjugate Gradient trainscg

Performance: Cross Entropy crossentropy

Calculations: MEX

Training Plots

[Performance](#)[Training State](#)[Error Histogram](#)[Confusion](#)[Receiver Operating Characteristic](#)

Training Confusion Matrix

Output Class	Target Class	
	1	2
1	8 57.1%	0 0.0%
2	0 0.0%	6 42.9%
	100% 0.0%	100% 0.0%

Validation Confusion Matrix

Output Class	Target Class	
	1	2
1	3 75.0%	0 0.0%
2	0 0.0%	1 25.0%
	100% 0.0%	100% 0.0%

Test Confusion Matrix

Output Class	Target Class	
	1	2
1	1 50.0%	0 0.0%
2	0 0.0%	1 50.0%
	100% 0.0%	100% 0.0%

All Confusion Matrix

Output Class	Target Class	
	1	2
1	12 60.0%	0 0.0%
2	0 0.0%	8 40.0%
	100% 0.0%	100% 0.0%

3. Consider the data in Table 1: [‘dataforquestion03assignment01.csv’](#) (uploaded in Avenue’s Assignment 01 folder).

(a) Use the linear programming classification method discussed in class to develop a neural network classifier that can be used for training the data in Table 1. Comment on the regions formed by the supermasks and foreign-object masks.

This data is highly non-linear. Therefore, we could not classify data using linear programming classification properly. As we can see, it could only predict data with 50% accuracy. Supermask could only classify one set. As there was no foreign object or no overlap in classes, foreign object class could not classify properly.

Neural Network Design and Training with LP

```
In 866 1 from pulp import *
```

```
In 867 1 prob = LpProblem("Neural Network Design and Training with LP", LpMinimize)
      2 probf = LpProblem("Neural Network Design and Training with LP", LpMinimize)
```

✓ C:\Users\malhoa25\Desktop\Q775 Machine Learning\venv\lib\site-packages\pulp\...
the name. Converted to '_'
warnings.warn("Spaces are not permitted in the name. Converted to '_'")

LP Supermask for Class A

```
In 868 1 # Variables
      2 rho1 = LpVariable("rho1", 0)
      3 rho2 = LpVariable("rho2", 0)
      4 rho3 = LpVariable("rho3", 0)
      5 rho4 = LpVariable("rho4", 0)
      6 rho5 = LpVariable("rho5", 0)
      7 rho6 = LpVariable("rho6", 0)
```

```

8 rho7 = LpVariable("rho7", 0)
9 rho8 = LpVariable("rho8", 0)
10 rho9 = LpVariable("rho9", 0)
11 rho10 = LpVariable("rho10", 0)
12 v1 = LpVariable("v1")
13 v2 = LpVariable("v2")
14 v3 = LpVariable("v3")
15 v4 = LpVariable("v4")
16 v5 = LpVariable("v5")
17 v6 = LpVariable("v6")

```

```

In 869 1 # Objective
2 prob += rho1+rho2+rho3+rho4+rho5+rho6+rho7+rho8+rho9+rho10

```

```

In 870 1 # Constraints
2 prob += 0.04*v1 +1.00*v2 + 0.20*v3 + 0.20*v4 +1.00*v5 + v6 - rho1 ==0
3 prob += 0.04*v1 +0.64*v2 + 0.16*v3 + 0.20*v4 +0.80*v5 + v6 - rho2 ==0
4 prob += 0.09*v1 +0.81*v2 + 0.27*v3 + 0.30*v4 +0.90*v5 + v6 - rho3 ==0
5 prob += 0.16*v1 +0.64*v2 + 0.32*v3 + 0.40*v4 +0.80*v5 + v6 - rho4 ==0
6 prob += 0.16*v1 +1.00*v2 + 0.40*v3 + 0.40*v4 +1.00*v5 + v6 - rho5 ==0
7 prob += 0.64*v1 +0.16*v2 + 0.32*v3 + 0.80*v4 +0.40*v5 + v6 - rho6 ==0
8 prob += 0.64*v1 +0.04*v2 + 0.16*v3 + 0.80*v4 +0.20*v5 + v6 - rho7 ==0
9 prob += 0.81*v1 +0.09*v2 + 0.27*v3 + 0.90*v4 +0.30*v5 + v6 - rho8 ==0
10 prob += 1.00*v1 +0.04*v2 + 0.20*v3 + 1.00*v4 +0.20*v5 + v6 - rho9 ==0
11 prob += 1.00*v1 +0.16*v2 + 0.40*v3 + 1.00*v4 +0.40*v5 + v6 - rho10 ==0
12 prob += v1 + v2 == -1

```

```
In 871 1 prob.solve()
```

```
Out 871 1
```

```
In 872 1 # The status of the solution is printed to the screen
      2 print("Status:", LpStatus[prob.status])
```

```
Status: Optimal
```

Solution

```
In 873 1 # Each of the variables is printed with it's resolved optimum value
      2 for v in prob.variables():
      3     print(v.name, "=", v.varValue)
```

```

      4
      5 rho5 = 0.0
      6 rho6 = 0.03
      7 rho7 = 0.0
      8 rho8 = 0.0175
      9 rho9 = 0.0
     10 v1 = -0.5
     11 v2 = -0.5
```

```
v3 = -0.75
v4 = 1.05
v5 = 1.05
v6 = -0.59
```

Normalized

```
In 874 1 from numpy import linalg as LA
```

```
In 875 1 vs=[]
      2 for v in range(10,16):
      3     vs.append(prob.variables()[v].varValue)
      4 vs
```

```
Out 875      [-0.5, -0.5, -0.75, 1.05, 1.05, -0.59]
```

```
In 876 1 vsnorm=[]
      2 for v in range(10,16):
      3     vsnorm.append(prob.variables()[v].varValue/LA.norm(vs))
      4 vsnorm
```

```
Out 876 1 [-0.26295402051409067,
      2      -0.26295402051409067,
```



```
.....,
-0.39443103077113606,
0.5522034430795905,
0.5522034430795905,
-0.310285744206627]
```

Foreign Object Mask for Class A

```
In 877 1  # Variables
2  eta1 = LpVariable("eta1", 0)
3  eta2 = LpVariable("eta2", 0)
4  eta3 = LpVariable("eta3", 0)
5  rhof1 = LpVariable("rhof1", 0)
6  rhof2 = LpVariable("rhof2", 0)
7  rhof3 = LpVariable("rhof3", 0)
8
9  vf1 = LpVariable("vf1")
10 vf2 = LpVariable("vf2")
11 vf3 = LpVariable("vf3")
12 vf4 = LpVariable("vf4")
13 vf5 = LpVariable("vf5")
14 vf6 = LpVariable("vf6")
```

```
In 878 1  # Objective
2  probf += eta2 #+ eta2 + eta3
```

```

In 879 1
2 # Constraints
3 probf += 0.04 * vf1 + 1.00 * vf2 + 0.2 * vf3 + 0.20 * vf4 + 1.00 * vf5 + vf6 <= 0
4 probf += 0.04 * vf1 + 0.64 * vf2 + 0.16 * vf3 + 0.20 * vf4 + 0.80 * vf5 + vf6 <= 0
5 probf += 0.09 * vf1 + 0.81 * vf2 + 0.27 * vf3 + 0.30 * vf4 + 0.90 * vf5 + vf6 <= 0
6 probf += 0.16 * vf1 + 0.64 * vf2 + 0.32 * vf3 + 0.40 * vf4 + 0.80 * vf5 + vf6 <= 0
7 probf += 0.16 * vf1 + 1.00 * vf2 + 0.40 * vf3 + 0.40 * vf4 + 1.00 * vf5 + vf6 <= 0
8 probf += 0.64 * vf1 + 0.16 * vf2 + 0.32 * vf3 + 0.80 * vf4 + 0.40 * vf5 + vf6 <= 0
9 probf += 0.64 * vf1 + 0.04 * vf2 + 0.16 * vf3 + 0.80 * vf4 + 0.20 * vf5 + vf6 <= 0
10 probf += 0.81 * vf1 + 0.09 * vf2 + 0.27 * vf3 + 0.90 * vf4 + 0.30 * vf5 + vf6 <= 0
11 probf += 1.00 * vf1 + 0.04 * vf2 + 0.20 * vf3 + 1.00 * vf4 + 0.20 * vf5 + vf6 <= 0
12 probf += 1.00 * vf1 + 0.16 * vf2 + 0.40 * vf3 + 1.00 * vf4 + 0.40 * vf5 + vf6 <= 0
13 probf += 0.04 * vf1 + 0.04 * vf2 + 0.04 * vf3 + 0.20 * vf4 + 0.20 * vf5 + eta1 - rhof1 == 0
14 probf += 0.04 * vf1 + 0.16 * vf2 + 0.08 * vf3 + 0.20 * vf4 + 0.40 * vf5 + eta2 - rhof2 == 0
15 probf += 0.09 * vf1 + 0.09 * vf2 + 0.09 * vf3 + 0.30 * vf4 + 0.30 * vf5 + eta3 - rhof3 == 0
16 probf += rhof2 + rhof2 + rhof3 == 1
17

```

```

In 880 1 probf.solve()

```

```

Out 880 1

```

```

In 881 1 # The status of the solution is printed to the screen
2 print("Status_foreign:", LpStatus[probf.status])

```



```
In 882 1 ##### Solution
2 # Each of the variables is printed with it's resolved optimum value
3 for vf in probf.variables():
4     print(vf.name, "=", vf.varValue)
```

```
eta2 = 0.0
eta3 = 1.0
rhof1 = 0.0
rhof2 = 0.0
rhof3 = 1.0
vf1 = 0.0
vf2 = 0.0
vf3 = 0.0
vf4 = 0.0
vf5 = 0.0
vf6 = 0.0
```

```
In 883 1 from numpy import linalg as LA
```

```
In 884 1 vsf = []
2 for vf in range(2, 8):
3     vsf.append(probf.variables()[vf].varValue)
```

```

2 vsf
3 vsfnorm = []
4 for vf in range(2, 8):
5     vsfnorm.append(probf.variables()[vf].varValue / LA.norm(vsf))
6 vsfnorm

```

Out 108 [0.0, 0.0, 0.08454898264477458, 0.0, -0.7045748610097204, 0.7045748610097204]

Plotting Supermask and Foreign Object Mask Class A

```

In 109 1 import seaborn as sns
2 %matplotlib inline
3 import numpy as np
4 import matplotlib.pyplot as plt
5 def axes():
6     plt.axhline(0, alpha=.1)
7     plt.axvline(0, alpha=.1)
8
9 x, y = np.linspace(0,1.1,100), np.linspace(0,1.1,100)
10 X, Y = np.meshgrid(x,y)
11 Z1 = vsnorm[0]*X**2 + vsnorm[1]*Y**2 + vsnorm[2]*X*Y + vsnorm[3]*X + vsnorm[4]*Y + vsnorm[5]
12 Z2 = vsfnorm[0]*X**2 + vsfnorm[1]*Y**2 + vsfnorm[2]*X*Y + vsfnorm[3]*X + vsfnorm[4]*Y + vsfnorm[5]
13 plt.figure()
14 cp1 = plt.contour(X, Y, Z1,[0],colors='g')

```


```

15 cp2 = plt.contour(X, Y, Z2,[0],colors='r')
16 score1=[]
17 score2=[]
18 color=[]
19 import pandas as pd
20 data = pd.read_csv("./data_for_question_03_assignment_01.csv", index_col=0)
21 for i in range(1,21):
22     score1.append(data.Score1[i])
23     score2.append(data.Score2[i])
24     if i < 11:
25         color.append(1)
26     else:
27         color.append(2)
28 import matplotlib.colors as cl
29 cmap = cl.LinearSegmentedColormap.from_list("", ["cornflowerblue","orange"])
30 plt.scatter(score1,score2,c=color, cmap=cmap)

```

Out 109 <matplotlib.collections.PathCollection at 0x1e75f3de400>

✓



(b) Graph (plot) the data set and comment on the appropriateness of the use of linear discriminant function for the separation of the two classes.

Given data is non linear and complex data, therefore it is inappropriate to use linear discriminant function. As this will not separate the classes properly.

Linear Discriminant Analysis

```
In 13 1 import pandas as pd
      2 import numpy as np
```

```
In 14 1 data = pd.read_csv("./data_for_question_03_assignment_01.csv", index_col=0)
```

```
In 15 1 data.head()
```

Out 15

< < 5 rows > > 5 rows × 3 columns				
Object	Score1	Score2	Class	
1	0.2	1.0	A	
2	0.2	0.8	A	
3	0.3	0.9	A	
4	0.4	0.8	A	
5	0.4	1.0	A	

```
In 17 1 from sklearn.discriminant_analysis import LinearDiscriminantAnalysis
```

```
In 18 1 X = data[['Score1', 'Score2']]
```

```
In 19 1 X.head()
```

Out 19 ▾

|< < 5 rows ▾ > >| 5 rows × 2 columns

Object ↕	Score1 ↕	Score2 ↕
1	0.2	1.0
2	0.2	0.8
3	0.3	0.9
4	0.4	0.8
5	0.4	1.0

```
In 20 1 print(type(X))
```

```
2 print(X.shape)
```

```
<class 'pandas.core.frame.DataFrame'>
(20, 2)
```

```
In 21 1 y = data.Class
```

```
2 y.head()
```

Out 21

5 rows Length: 5, dtype: object

Object	Class
1	A
2	A
3	A
4	A
5	A

```
In 22 1 print(type(y))
      2 print(y.shape)
```

```
<class 'pandas.core.series.Series'>
(20,)
```

```
In 23 1 import seaborn as sns
      2 import matplotlib.pyplot as plt
      3 %matplotlib inline
```


```
In 24 1 from sklearn.discriminant_analysis import LinearDiscriminantAnalysis
      2
      3 lda = LinearDiscriminantAnalysis()
      4 lda.fit(X, y)
      5
      6 coef = lda.coef_[0]
      7 intercept = lda.intercept_
      8
```



```

10 import matplotlib.pyplot as plt
11
12 sns.lmplot(x='Score1', y='Score2', data=data, hue='Class', fit_reg=False)
13
14 x_vals = np.array([0, 1])
15 y_vals = -(coef[0] * x_vals + intercept) / coef[1]
16
17 plt.plot(x_vals, y_vals, 'r--', linewidth=2)
18
19 plt.show()
20

```

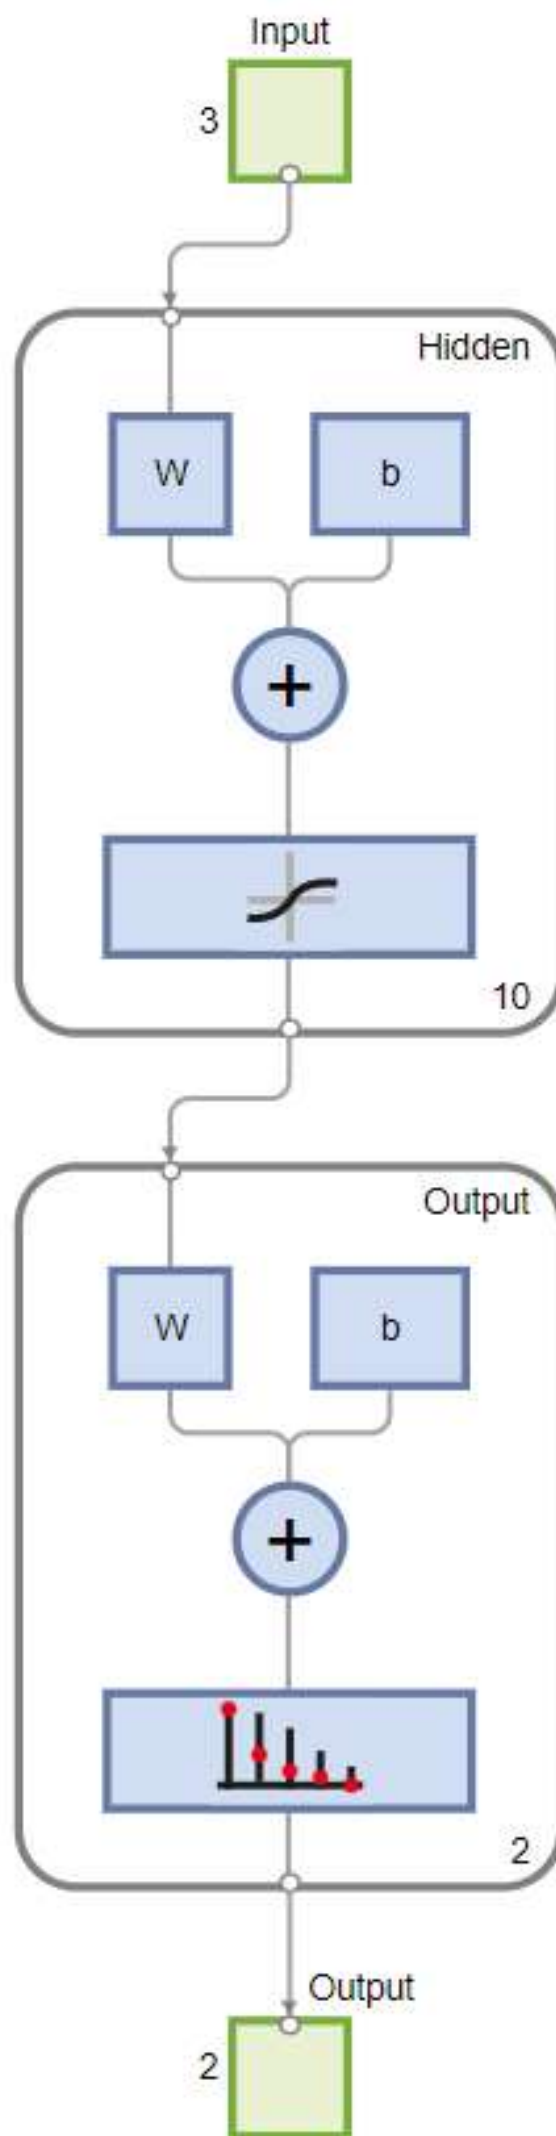


(c) Use the appropriate Neural Network code (MATLAB or Python) to obtain a neural network classifier for the same data set. Compare this result to that obtained in part 3(a).

Feedforward artificial neural network is used to predict the classification of the data. Referring to the confusion matrix, it shows that it has predicted the data with 100% accuracy. But when classification is done using linear programming or linear discrimination method, these methods could not classify the objects with proper boundaries.

clear; clc;	
%% Loading the data_for_question_03_assignment_01 data, with the classes at the last column	
A = load('ans3c_data.txt');	
% C is the matrix of the feature vectors C = A(1:end,1:end-1)';	
%% Create matrix for classes (targets) T=zeros(size(unique(A(:, end)),1),size(A,1));	
for i=1:size(A,1) T(A(i,end),i)=1; %create identity vectors to reflect class type end	
%% Initializing a neural network 'net' net = feedforwardnet; net = configure(net, C, T); % net - NN, C-feature vector, T-target	
%% Setting hidden layer hiddenLayerSize = 10; % Setting the number of hidden layers to 10 net = patternnet(hiddenLayerSize); % Pattern recognition network	
%% Split data for training, validation and testing net.divideParam.trainRatio = 0.7; % Ratio of training data is 70% net.divideParam.valRatio = 0.2; % Ratio of validation data is 20% net.divideParam.testRatio = 0.1; % Ratio of testing data is 10%	
%% Training the network [net, tr] = train(net, C, T); % the resulting model is the output net	
outputs = net(C);	
%% Computing the classification errors errors = gsubtract(T, outputs); performance = perform(net, T, outputs); view(net)	

Pattern Recognition Neural Network (view)



[Network Diagram](#)

Training Results

Training finished: Reached minimum gradient ✓

Training Progress

Unit	Initial Value	Stopped Value	Target Value	
Epoch	0	25	1000	▲
Elapsed Time	-	00:00:00	-	
Performance	0.828	3.03e-07	0	
Gradient	0.593	6.95e-07	1e-06	
Validation Checks	0	0	6	▼

Training Algorithms

Data Division: Random dividerand

Training: Scaled Conjugate Gradient trainscg

Performance: Cross Entropy crossentropy

Calculations: MEX

Training Plots

[Performance](#)[Training State](#)[Error Histogram](#)[Confusion](#)[Receiver Operating Characteristic](#)

Training Confusion Matrix

Output Class	1	2	
	7 50.0%	0 0.0%	100% 0.0%
	0 0.0%	7 50.0%	100% 0.0%
	100% 0.0%	100% 0.0%	100% 0.0%
	↘	↗	
	Target Class		

Validation Confusion Matrix

Output Class	1	2	
	2 50.0%	0 0.0%	100% 0.0%
	0 0.0%	2 50.0%	100% 0.0%
	100% 0.0%	100% 0.0%	100% 0.0%
	↘	↗	
	Target Class		

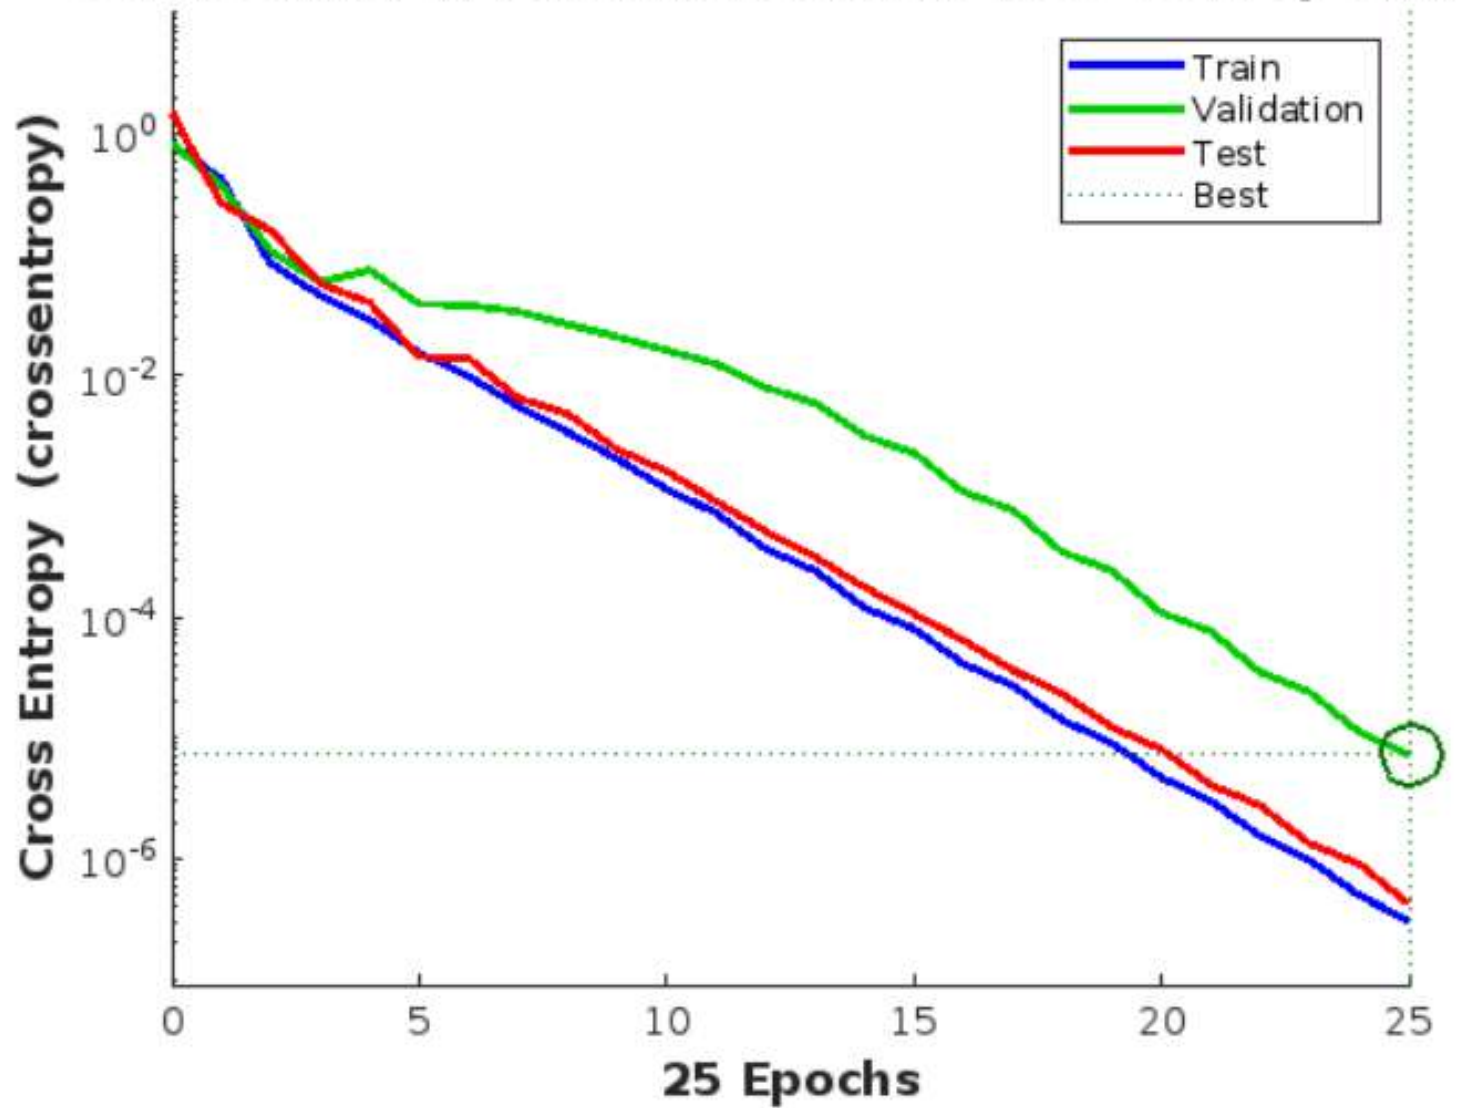
Test Confusion Matrix

Output Class	1	2	
	1 50.0%	0 0.0%	100% 0.0%
	0 0.0%	1 50.0%	100% 0.0%
	100% 0.0%	100% 0.0%	100% 0.0%
	↘	↗	
	Target Class		

All Confusion Matrix

Output Class	1	2	
	10 50.0%	0 0.0%	100% 0.0%
	0 0.0%	10 50.0%	100% 0.0%
	100% 0.0%	100% 0.0%	100% 0.0%
	↘	↗	
	Target Class		

Best Validation Performance is 7.2465e-06 at epoch 25



(d) Assuming that we do not know the data classes and assignments, using the graph in part 3(b) how many clusters would you suspect existed and for which objects? What is the basis of your clustering?

When we do not know the data classes and assignments, then we can use elbow method which is a graphical tool to estimate clusters. Using kmeans model, elbow curve is formed. And we can see, the elbow is formed at $k = 2$, showing that given data is clustered in total two groups.

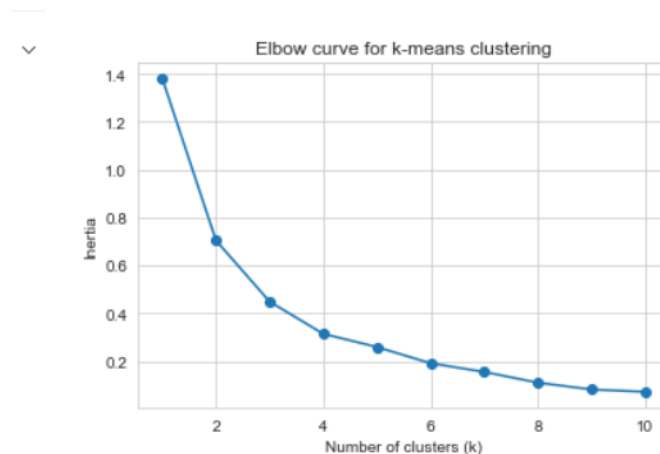


Figure 1: Elbow Curve showing $k=2$

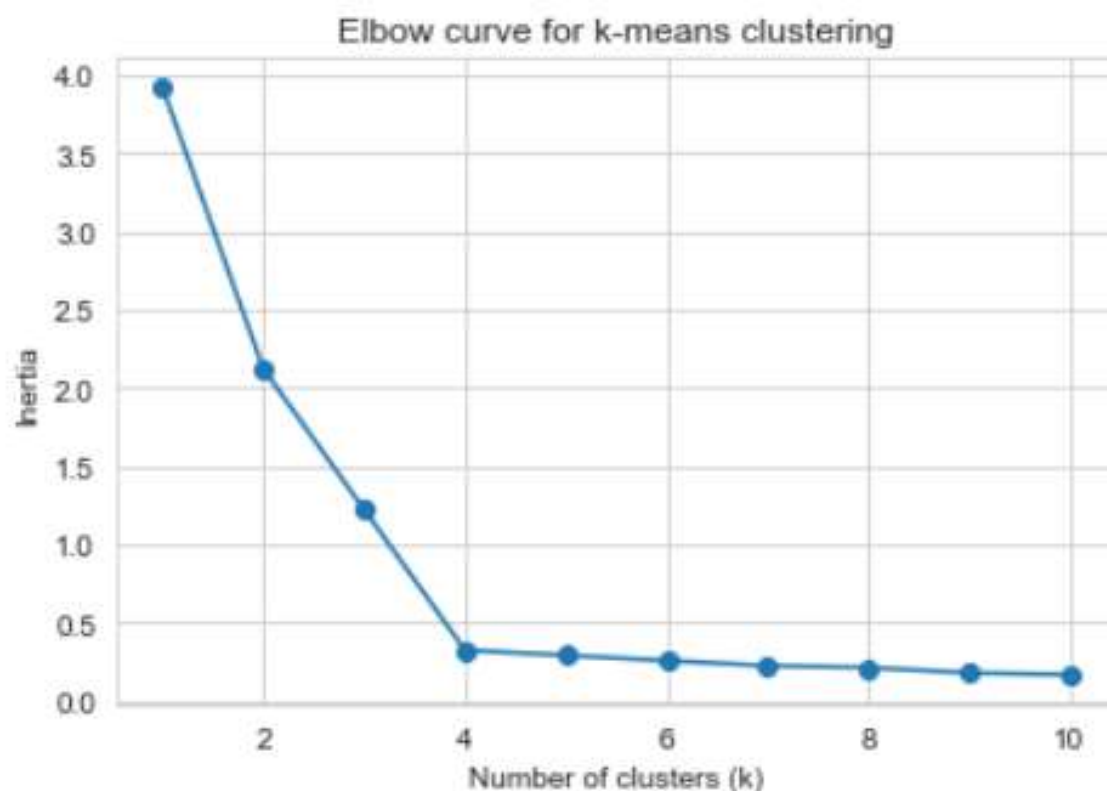
(e) Develop a clustering model and use it to find the optimal number of clusters and their objects.

Using elbow method, optimal number of clusters and objects are formed.


```
In 3 1 import pandas as pd
      2 import numpy as np
      3 from sklearn.cluster import KMeans
      4 import matplotlib.pyplot as plt
      5
      6 # read data from CSV file
      7 data = pd.read_csv('data_for_question_03_assignment_01.csv', index_col='Object')
      8
      9 # drop the Class column as it's not needed for clustering
     10 X = data.drop('Class', axis=1)
     11
     12 # perform clustering for different values of k`
     13 inertias = []
     14 k_values = range(1, 11)
     15 for k in k_values:
     16     model = KMeans(n_clusters=k)
     17     model.fit(X)
     18     inertias.append(model.inertia_)
     19
     20 # plot the elbow curve to find the optimal number of clusters
     21 plt.plot(k_values, inertias, '-o')
     22 plt.xlabel('Number of clusters (k)')
     23 plt.ylabel('Inertia')
     24 plt.title('Elbow curve for k-means clustering')
     25 plt.show()
     26
```

```
warnings.warnl
```

▼



```
In 4 1 # fit the KMeans model with k=2
      2 model = KMeans(n_clusters=2)
      3 model.fit(X)
      4
      5 # print the cluster labels
      6 labels = pd.Series(model.labels_, index=X.index)
      7 print(labels)
```

▼

Object

1	1
2	1
3	1
4	1
5	1
6	0
7	0
8	0
9	0
10	0
11	1
12	1
13	1
14	1
15	1
16	0
17	0
18	0
19	0
20	0

dtype: int32

4. In signed network/graph, each edge (link) is labelled as either positive or negative to represent similarity or dissimilarity between connecting nodes (vertices), respectively. Correlation Clustering (CC) problem (Bansal et al., 2004) is a well-known graph/network partitioning problem that finds clusters in given signed network. The integer linear programming programming (ILP) formulation of the CC problem for a general weighted signed network $G(V, E)$ can be defined as follows (Demaine et al., 2006):

$$\begin{aligned}
\min \quad & \sum_{(i,j) \in E^+} w_{ij} x_{ij} + \sum_{(i,j) \in E^-} w_{ij} (1 - x_{ij}) \\
\text{s.t.} \quad & x_{ij} + x_{jk} \geq x_{ik} \quad ; \forall i, j, k \in V \\
& x_{ij} = x_{jk}; \quad \forall i, j \in V \\
& x_{ij} \in \{0, 1\}; \quad \forall i, j \in V
\end{aligned}$$

(a) Implement the above ILP formulation using any language in your choice and package (not in excel).

(b) Use your ILP implementation to identify clusters in the Slovene Parliamentary Political party's mutual relationship data set '[slovenepoliticalparties.csv](#)' (Avenue's Assignment 01 folder). The description of this data can be found in here:

ILP formulation is done in Python and part2 a and b answer is shown below.

<http://networkdata.ics.uci.edu/netdata/html/Stranke94.html> .

```
In 121 1 from docplex.mp.model import Model
      2 import sys
      3 sys.path.insert(1, './network_importer')
      4 import network_importer as nimp
```

Read Network

```
In 122 1 G = nimp.networkx_read_weighted_network_from_csv('./test_net_assign1.csv')
```

```
In 123 1 print(G.nodes)
```

```
[1, 2, 3, 4, 5, 6, 7, 8, 9, 10]
```

```
In 124 1 print(G.edges)
```

```
[(1, 2), (1, 3), (1, 4), (1, 5), (1, 6), (1, 7), (1, 8), (1, 9), (1, 10), (2, 3), (2, 4), (2, 5), (2, 6), (2, 7), (2, 8), (2, 9), (2, 10), (3, 4), (3, 5), (3, 6), (3, 7), (3, 8), (3, 9), (3, 10), (4, 5), (4, 6), (4, 7), (4, 8), (4, 9), (4, 10), (5, 6), (5, 7), (5, 8), (5, 9), (5, 10), (6, 7), (6, 8), (6, 9), (6, 10), (7, 8), (7, 9), (7, 10), (8, 9), (8, 10), (9, 10)]
```

ILP Model : Decision variables

```
In 125 1  # Index range
2  N = range(1, len(G.nodes) + 1)
3
4  # Model
5  ce = Model(name='clustering_editing')
6
7  # Decision variable
8  x = {(i, j): ce.binary_var(name='x_{0}_{1}'.format(i, j)) for i in N for j in N}
```

ILP Model: Objective Function

```
In 126 1  # Objective function
2  LDC: int = 0          # total edge deletion cost
3  LIC: int = 0          # total edge insertion cost
4  for i in N:
5      for j in N:
6          if i < j:
7              if G.has_edge(i, j)>0:
8                  w_ij = G.edges[i, j]['weight']
9
10         LDC_this = ce.sum(x[i, j] * w_ij)
```

```

11 LDC = LDC + LDC_this
12 else:
13     del_ij = 1
14
15     LIC_this = ce.sum((1 - x[i, j]) * del_ij)
16     LIC = LIC + LIC_this

```

```

In 127 1 # Total Cost (Objective function) = Link Deletion Cost (LDC) + Link Insertion Cost (LIC)
2     total_cost = LDC + LIC
3
4     # Minimize all cost
5     ce.minimize(total_cost)

```

Constraint

```

In 128 1 # Constrain: Triangle inequality
2     for i in N:
3         for j in N:
4             for k in N:
5                 if i != j != k:
6                     # linear expression for constraint
7                     c1_this = ce.sum(x[i, j] + x[j, k] - x[i, k])
8
9                     # set constraint
10                    ce.add_constraint(c1_this >= 0)
11
12 # Constraint: Undirected link

```

```

12 # Constraint: Undirected link
13 for i in N:
14     for j in N:
15         if i != j != k:
16             # linear expression for constraint
17             c2_this = ce.sum(x[i, j] - x[j, k])
18
19             # set constraint
20             ce.add_constraint(c2_this == 0)

```

Solve Model

```

In 129 1 # Solve model
2       solution = ce.solve()
3
4       if solution is None:
5           print('- Model is infeasible')
6       else:
7           print('Model found an optimal solution')
8
9       assert solution

```

Model found an optimal solution


```
In 130 1 print('Objective value: ', solution.objective_value)
```

Objective value: -1881.0

Returning Clusters

```
In 131 1 sys.path.insert(1, './partitioning_for_ce')
      2 import partitioning_for_ce as part
```

```
In 132 1 # Get clusters
      2 P = [] # clusters partition set
      3
      4 for i in N:
      5     for j in N:
      6         if i < j:
      7             if solution.get_value(x[i, j]) >= 0:
      8                 S = []
      9                 S.append(i)
     10                 S.append(j)
     11                 P = part.get_partitions(P, S)
     12
     13 # Merging any two partitions that have common element(s)
     14 P_final: list = []
```

```
15 for L in P:  
16     P_final = part.get_partitions(P_final, L)  
17
```

```
In 133 1 print("Final Clusters:" , P_final)
```

```
Final Clusters: [[1, 2, 3, 4, 5, 6, 7, 8, 9, 10]]
```

The given model second constraint should be $x_{ji} = x_{ij}$. So I have re-coded to get better results and form clusters.

Clustering Editing

```
In 46 1 import ...  
      2  
      3  
      4 sys.path.insert(1, './network_importer')  
      5 import network_importer as nimp  
      6  
      7
```

Read Network

```
In 47 1 ### Read Network  
      2 G = nimp.networkx_read_weighted_network_from_csv('./test_net_assgn1.csv')
```

```
In 48 1 print(G.nodes)  
  
      [1, 2, 3, 4, 5, 6, 7, 8, 9, 10]
```

```
In 49 1 print(G.edges)
```

```
[1, 2, 3, 4, 5, 6, 7, 8, 9, 10]
```

```
In 49 1 print(G.edges)
```

```
✓ [(1, 2), (1, 3), (1, 4), (1, 5), (1, 6), (1, 7), (1, 8), (1, 9), (1, 10), (2, 3), (2, 4), (2, 5), (2, 6), (2, 7), (2, 8), (2, 9), (2, 10), (3, 4), (3, 5), (3, 6), (3, 7), (3, 8), (3, 9), (3, 10), (4, 5), (4, 6), (4, 7), (4, 8), (4, 9), (4, 10), (5, 6), (5, 7), (5, 8), (5, 9), (5, 10), (6, 7), (6, 8), (6, 9), (6, 10), (7, 8), (7, 9), (7, 10), (8, 9), (8, 10), (9, 10)]
```

ILP Model : Decision variables

```
In 50 1 # Index range
2 N = range(1, len(G.nodes) + 1)
3
4 # Model
5 ce = Model(name='clustering_editing')
6
7 # Decision variable
8 x = {(i, j): ce.binary_var(name='x_{0}_{1}'.format(i, j)) for i in N for j in N}
```

ILP Model: Objective Function

```

In 51  1  # Objective function
      2  LDC: int = 0          # total edge deletion cost
      3  LIC: int = 0          # total edge insertion cost
      4  for i in N:
      5      for j in N:
      6          if i < j:
      7              if G.has_edge(i, j)>0:
      8                  w_ij = G.edges[i, j]['weight']
      9
     10                  LDC_this = ce.sum(x[i, j] * w_ij)
     11                  LDC = LDC + LDC_this
     12              else:
     13                  del_ij = 1
     14
     15                  LIC_this = ce.sum((1 - x[i, j]) * del_ij)
     16                  LIC = LIC + LIC_this

```

```

In 52  1  # Total Cost (Objective function) = Link Deletion Cost (LDC) + Link Insertion Cost (LIC)
      2  total_cost = LDC + LIC
      3
      4  # Minimize all cost
      5  ce.minimize(total_cost)

```

Constraint

```

In 53 1  # Constrain: Triangle inequality
2  for i in N:
3      for j in N:
4          for k in N:
5              if i != j != k:
6                  # linear expression for constraint
7                  c1_this = ce.sum(x[i, j] + x[j, k] - x[i, k])
8
9                  # set constraint
10             ce.add_constraint(c1_this >= 0)
11
12  # Constraint: Undirected link
13  for i in N:
14      for j in N:
15          if i != j:
16              # linear expression for constraint
17              c2_this = ce.sum(x[i, j] - x[j, i])
18
19              # set constraint
20             ce.add_constraint(c2_this == 0)

```

Solve Model

```
In 54 1  # Solve model
      2  solution = ce.solve()
      3
      4  if solution is None:
      5      print('- Model is infeasible')
      6  else:
      7      print('Model found an optimal solution')
      8
      9  assert solution
```

Model found an optimal solution

```
In 55 1  print('Objective value: ', solution.objective_value)
```

Objective value: -4193.0

Returning Clusters

[Add Code Cell](#)[Add Markdown Cell](#)

```
In 56 1 sys.path.insert(1, './week05/partitioning_for_ce')
      2 import partitioning_for_ce as part
```

```
In 57 1 # Get clusters
      2 P = [] # clusters partition set
      3
      4 for i in N:
      5     for j in N:
      6         if i < j:
      7             if solution.get_value(x[i, j]) == 0:
      8                 S = []
      9                 S.append(i)
     10                 S.append(j)
     11                 P = part.get_partitions(P, S)
     12
     13 # Merging any two partitions that have common element(s)
     14 P_final: list = []
     15 for L in P:
     16     P_final = part.get_partitions(P_final, L)
     17
```

```
In 58 1 print("Final Clusters:" , P_final)
```

```
Final Clusters: [[2, 4, 5, 7, 10], [1, 3, 6, 8, 9]]
```

(c) Compare and discuss your result with the clusters reported in Figure 10 in Wu et al. (2016). URL: <https://www.sciencedirect.com/science/article/abs/pii/S037843711500802X>

Formulation 1: As per given in homework ILP formulation given in question is implemented in Python and only one cluster $[1, 2, 3, 4, 5, 6, 7, 8, 9, 10]$ is formed. However in Figure 10 in Wu et al. (2016) clusters $[2, 4, 5, 7, 10]$ and $[1, 3, 6, 8, 9]$ are formed. Even though signed link network is formed in both the ILP formulation but in paper second constraint is $x_{ij} = x_{ji}$ instead of $x_{ij} = x_{jk}$ which is given in the question. To obtain clusters for the given question we should be either changing the objective or to add more constraint as a penalty if no clusters are formed. Currently, objective is just minimizing the cost of cluster editing and is not asking to form number of clusters. So its finding an optimal result with single cluster only.

Formulation 2: Where second constraint is corrected Cluster is obtained similar to the given in paper.

5. For this question, use the BigQuery public data '[bigquery.public data.new.york.tlc.yellow.trips.2012](#)' on the Google Cloud Platform

(a) Use the 'RAND()' to pull 1 in 100,000 records. This will result in approximately 1,793 records (1 in 100,000 records). Print the first 10 entries of the extracted data.

```

In 1372 1 CREDENTIAL_PATH = './q775-379622-e6f9f62f9691.json'
        2 credentials = service_account.Credentials.from_service_account_file(CREDENTIAL_PATH,)
        3 project_id = 'q775-379622'

        ## Querying New York data from BigQuery

In 1373 1 sql_query = \
        2 '''SELECT *
        3 FROM bigquery-public-data.new_york.tlc_yellow_trips_2012
        4 where rand() < 1/100000;
        5 '''

In 1374 1 trips = pd.read_gbq(sql_query, project_id=project_id, credentials=credentials)

In 1375 1 trips.shape

Out 1375 (1833, 19)

```

Figure 2: Used SQL query to pull 1 in 100,000 records.'

I got 1833 records and 19 rows.

(b) Visualize the trip distance versus fare amount. Do you observe any anomalies? If yes, clean the data by removing the invalid data and running a new BigQuery or SQL query. Then plot the data again.


Here we see few outliers. Now we will draw box plot to understand its quartile range and then remove them and then finally drawing scatter plot again to check if all the outliers have been removed. Finally in last figure all the outliers has been removed.

```
In 1377 1 # Read dataframe from local CSV file
2 trips['trip_distance'] = trips['trip_distance'].astype(float)
3 trips['fare_amount'] = trips['fare_amount'].astype(float)
```

```
In 1378 1 sns.displot(trips['trip_distance']).set(title='trip_distance')
2
```

Out 1378 <seaborn.axisgrid.FacetGrid at 0x1dfcc2833d0>


✓



Add Code


```
In 1379 1 sns.displot(trips['fare_amount']).set(title='fare_amount')
```

```
Out 1379 <seaborn.axisgrid.FacetGrid at 0x1dfcc340550>
```



```
In 1380 1 sns.scatterplot(data=trips, x='trip_distance', y='fare_amount')
```

```
Out 1380 <AxesSubplot:xlabel='trip_distance', ylabel='fare_amount'>
```




```
In 1382 1  #Step-5: Finding the IQR
        2
        3 percentile25 = trips['trip_distance'].quantile(0.25)
        4 percentile75 = trips['trip_distance'].quantile(0.75)
        5
        6 print("75th quartile: ",percentile75)
        7 print("25th quartile: ",percentile25)

        75th quartile:  3.12
        25th quartile:  1.03
```

```
In 1383 1  iqr = percentile75 - percentile25
        2  print ("IQR: ",iqr)

        IQR:  2.09
```


In 1384


```
1
2 #Step-6: Finding the upper and lower limits
3
4 upper_limit = percentile75 + 1.5 * iqr
5 lower_limit = percentile25 - 1.5 * iqr
6 #Step-7: Finding outliers
7
8 trips[trips['trip_distance'] > upper_limit]
9 trips[trips['trip_distance'] < lower_limit]
10
11 #Step-8: Trimming outliers
12
13 trips[trips['trip_distance'] < upper_limit]
14 new_trips = trips[trips['trip_distance'] < upper_limit]
15 new_trips.shape
```

Out 1384

(1677, 19)


```
In 1385 1 sns.boxplot(data=trips['trip_distance'])
```

Out 1385 <AxesSubplot:>



```
In 1386 1      #Step-9: Compare the plots after trimming
        2      sns.boxplot(data=new_trips['trip_distance'])
```

Out 1386 <AxesSubplot:>




```
In 1387 1 sns.boxplot(trips['fare_amount'])
```

⌵ C:\Users\malhoa25\Anaconda3\lib\site-packages\se
arg: x. From version 0.12, the only valid position
keyword will result in an error or misinterpretation
warnings.warn(

```
Out 1387 <AxesSubplot:xlabel='fare_amount'>
```

⌵



```
In 1388 1  #Step-5: Finding the IQR
2
3  percentile25 = trips['fare_amount'].quantile(0.25)
4  percentile75 = trips['fare_amount'].quantile(0.75)
5
6  print("75th quartile: ",percentile75)
7  print("25th quartile: ",percentile25)
```

75th quartile: 12.1

25th quartile: 6.0

```
In 1389 1  iqr = percentile75 - percentile25
2  print ("IQR: ",iqr)
```


IQR: 6.1

```
In 1390 1 upper_limit = percentile75 + 1.5 * iqr
2 lower_limit = percentile25 - 1.5 * iqr
3 #Step-7: Finding outliers
4
5 trips[trips['fare_amount'] > upper_limit]
6 trips[trips['fare_amount'] < lower_limit]
7
8 #Step-8: Trimming outliers
9
10 trips[trips['fare_amount'] < upper_limit]
11 new_trips = trips[trips['fare_amount'] < upper_limit]
12 new_trips.shape
```

Out 1390 (1695, 19)

```
In 1391 1 sns.boxplot(data=new_trips['fare_amount'])
```

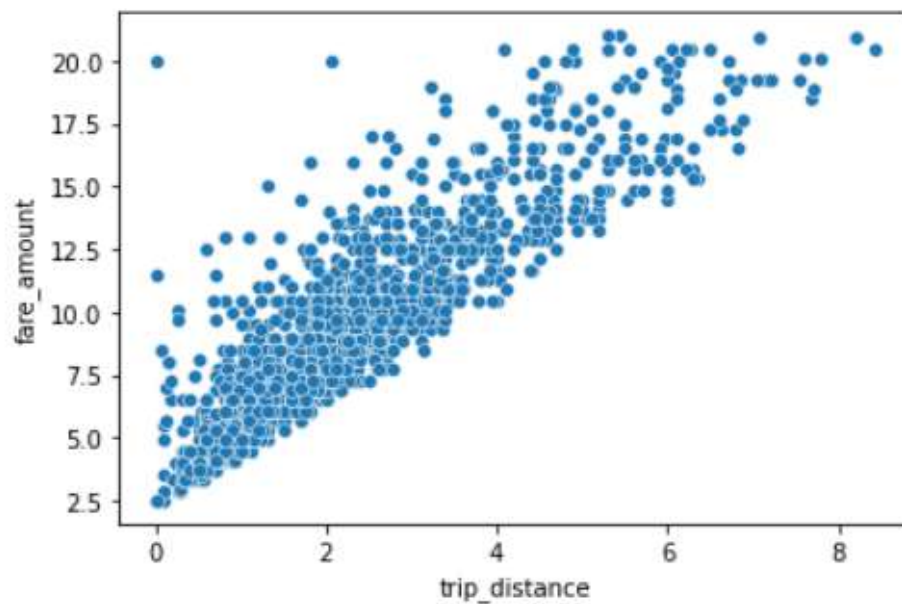
Out 1391 <AxesSubplot:>



```
In 1392 1 sns.scatterplot(data=new_trips, x='trip_distance', y='fare_amount')
```

```
Out 1392 <AxesSubplot:xlabel='trip_distance', ylabel='fare_amount'>
```

✓



```
In 1393 1 new_trips.shape
```

```
Out 1393 (1695, 19)
```


(c) Examine whether the toll amount is captured in the total amount by looking at records that have positive tolls for pick-up time 2012-09-05 15:45:00. What do you notice?

```
In 1394 1  sql_query2 = \
2  '''SELECT fare_amount, extra, mta_tax, tip_amount, tolls_amount, total_amount
3  FROM bigquery-public-data.new_york.tlc_yellow_trips_2012
4  where (pickup_datetime = '2012-09-05 15:45:00') and (tolls_amount != 0) ;
5  '''

In 1395 1  trips2 = pd.read_gbq(sql_query2, project_id=project_id, credentials=credentials)
2  trips2.shape

Out 1395  (13, 6)
```

```
In 1396 1  print(trips2)
```

	fare_amount	extra	mta_tax	tip_amount	tolls_amount	total_amount
0	38.0	0.0	0.5	0.0	4.8	43.3
1	32.0	0.0	0.5	9.6	4.8	46.9
2	68.5	0.0	0.0	0.0	9.5	78.0
3	32.5	0.0	0.5	0.0	4.8	37.8
4	52.0	0.0	0.5	10.4	4.8	67.7
5	27.0	0.0	0.5	5.4	4.8	37.7
6	28.0	0.0	0.5	5.6	4.8	38.9
7	34.0	0.0	0.5	6.8	4.8	46.1
8	32.5	0.0	0.5	0.0	4.8	37.8
9	53.0	0.0	0.5	10.6	4.8	68.9

Figure 3: Examining if toll amount is captured in the total amount

After examining, I found that toll amount is captured in the total amount. We can analyse in the figure.

(d) Report the descriptive statistics on the selected data: count, mean, std, min, 25%, 50%, 75% and max. What do you notice about the longitude values?

In 1397 1 new_trips.describe()

Out 1397

	passenger_count	trip_distance	pickup_longitude	pickup_latitude	rate_code	dropoff_longitude	dropoff_latitude
count	1695.000000	1695.000000	1695.000000	1695.000000	1695.000000	1695.000000	1695.000000
mean	1.719764	2.058490	-72.195847	39.790207	1.004720	-72.232519	39.791000
std	1.340472	1.453646	11.342457	6.187833	0.113863	11.232923	6.187900
min	0.000000	0.000000	-74.017884	0.000000	1.000000	-74.017936	0.000000
25%	1.000000	1.000000	-73.992848	40.735292	1.000000	-73.991480	40.735900
50%	1.000000	1.690000	-73.982555	40.751894	1.000000	-73.980987	40.753200
75%	2.000000	2.655000	-73.969731	40.766415	1.000000	-73.967162	40.768600
max	6.000000	8.430000	0.000000	40.847300	4.000000	0.000000	40.847300

Figure 4: Descriptive Statistics on the data

After looking at the descriptive data longitudinal values, I analyzed that even though this data has New York longitude i.e. values around -74 but few of the cells have discrepancy as maximum value shown is zero.

(e) Plot the start and end of a few (10) trips and their corresponding tolls. What do you notice?

After analyzing the scatter plot, we noticed there is no relationship between the distance travelled along with tolls amount. It may be happening that we have small data or may be the toll is in the city only. And if anyone taking the highway, needs to pay the toll.

```

In 10 1 sql_query3 = \
2      '''SELECT pickup_longitude, pickup_latitude, dropoff_latitude, dropoff_longitude, tolls_amount
3      FROM bigquery-public-data.new_york.tlc_yellow_trips_2012
4      where (tolls_amount != 0)
5      LIMIT 10;
6      '''
7  trips3 = pd.read_gbq(sql_query3, project_id=project_id, credentials=credentials)
8  trips3.shape

```

Out 10 (10, 5)

```

In 11 1 print(trips3)

```

	pickup_longitude	pickup_latitude	dropoff_latitude	dropoff_longitude	\
0	-73.776883	40.645851	40.795818	-73.962699	
1	-73.948316	40.829597	40.803214	-73.948467	
2	-73.976826	40.759201	40.751768	-73.991501	
3	-74.006153	40.748597	40.760968	-73.973491	
4	-73.984777	40.721880	40.732299	-74.044324	
5	-74.182947	40.628290	40.627548	-74.181620	
6	-74.039399	40.730083	40.730106	-74.039349	
7	-73.938656	40.752091	40.752091	-73.938656	
8	-73.982093	40.762020	40.737940	-74.059900	
9	-73.987108	40.769708	40.743072	-73.915164	

```

In 23 1 pickup_coords = trips3[['pickup_longitude', 'pickup_latitude']]
2      dropoff_coords = trips3[['dropoff_longitude', 'dropoff_latitude']]
3      tolls = trips3['tolls_amount']

```

```

In 30 1 plt.scatter(pickup_coords['pickup_longitude'], pickup_coords['pickup_latitude'], s=tolls*20, c='r', label='Pickup')
2      plt.scatter(dropoff_coords['dropoff_longitude'], dropoff_coords['dropoff_latitude'], s=tolls*20, c='b', label='Dropoff')
3
4      # Set the x and y limits to zoom in on a particular area
5      plt.xlim([-73.77, -74.2])
6      plt.ylim([40.6, 40.9])
7      # Plot the route for each trip
8      for i in range(len(trips3)):
9          pickup = (trips3['pickup_longitude'][i], trips3['pickup_latitude'][i])
10         dropoff = (trips3['dropoff_longitude'][i], trips3['dropoff_latitude'][i])
11         plt.plot([pickup[0], dropoff[0]], [pickup[1], dropoff[1]], 'k-', alpha=0.3)
12         plt.annotate('$'+str(trips3['tolls_amount'][i]), xy=(dropoff[0], dropoff[1]), xytext=(-20, 20), textcoords='offsetpoints',
13                    ha='right', va='bottom', bbox=dict(boxstyle='round,pad=0.5', fc='white', alpha=0.5))
14
15 # Display the plot
16 plt.show()

```



```
In 1401 1 new_trips = new_trips[new_trips['pickup_latitude'] != 0]
2 new_trips = new_trips[new_trips['pickup_latitude'] > 1]
3 new_trips = new_trips[new_trips['passenger_count'] != 0]
```

```
In 1402 1 new_trips.shape
```

```
Out 1402 (1647, 19)
```

```
In 1403 1 print(new_trips)
```

```

  -      - - - -
  2      NaN      17.50
  4      NaN      19.44
  5      NaN       9.38
  ...      ...      ...
1825     NaN      10.70
1826     NaN      12.24
1827     NaN      13.62
1831     NaN      11.70
1832     NaN       4.20
```

```
[1647 rows x 19 columns]
```

```

In 1404 1 #Replacing new values in total_amount by emptying it first and then finally adding fare and toll amount (partf)
2 new_trips.total_amount = "-"
3 new_trips.total_amount = new_trips.fare_amount + new_trips.tolls_amount

```

```

In 1405 1 new_trips.head()

```

Out 1405

5 rows x 19 columns								CSV				
	vendor_id	pickup_datetime	dropoff_datetime	passenger_count	trip_distance	pickup_longitude	pick					
0 CMT		2012-11-14 22:08:00+00:00	2012-11-14 23:20:28+00:00	1	3.90	-74.011093						
1 VTS		2012-09-07 12:22:00+00:00	2012-09-07 12:47:00+00:00	1	2.52	-73.963615						
2 VTS		2012-10-17 00:25:00+00:00	2012-10-17 00:45:00+00:00	4	3.76	-73.987095						
4 CMT		2012-07-24 07:32:18+00:00	2012-07-24 07:45:48+00:00	1	6.00	-74.016593						
5 VTS		2012-05-24 00:47:00+00:00	2012-05-24 00:53:00+00:00	1	1.75	-74.000945						

```

In 1406 1 #dropping column fare amount and toll amount
2 new_trips = new_trips.drop(['fare_amount'], axis=1)
3 new_trips

```

Out 1406

1547 rows x 10 columns										CSV			
longitude	dropoff_latitude	payment_type	extra	mta_tax	tip_amount	tolls_amount	imp_surcharge	total_amount					
-73.984400	40.691852	CRD	0.5	0.5	3.30	0.0	NaN	15.5					
-74.005605	40.745132	CSH	0.0	0.5	0.00	0.0	NaN	17.0					
-74.006200	40.706322	CSH	0.5	0.5	0.00	0.0	NaN	16.5					
-73.972623	40.754156	CRD	0.0	0.5	3.24	0.0	NaN	15.7					
-73.984243	40.743245	CRD	0.5	0.5	1.48	0.0	NaN	6.9					
...					
-74.000485	40.737380	CSH	0.5	0.5	0.00	0.0	NaN	9.7					

```
In 1407 1 new_trips = new_trips.drop(['tolls_amount'], axis=1)
2 new_trips = new_trips.drop(['rate_code'], axis=1)
3 new_trips = new_trips.drop(['store_and_fwd_flag'], axis=1)
4 new_trips = new_trips.drop(['payment_type'], axis=1)
5 new_trips = new_trips.drop(['imp_surcharge'], axis=1)
6 new_trips
```

Out 1407

11 rows x 13 columns

CSV

	vendor_id	pickup_datetime	dropoff_datetime	passenger_count	trip_distance	pickup_longitude	pic
2 VTS		2012-10-17 00:23:00+00:00	2012-10-17 00:43:00+00:00	1	3.70	-73.97073	
4 CMT		2012-07-24 07:32:18+00:00	2012-07-24 07:45:48+00:00	1	6.00	-74.016593	
5 VTS		2012-05-24 00:47:00+00:00	2012-05-24 00:53:00+00:00	1	1.75	-74.000945	
...	
1825 VTS		2012-05-31 01:55:00+00:00	2012-05-31 02:07:00+00:00	3	2.84	-73.970242	
1826 CMT		2012-03-05 10:37:12+00:00	2012-03-05 10:51:37+00:00	1	2.10	-73.979540	
1827 VTS		2012-09-22 18:07:00+00:00	2012-09-22 18:18:00+00:00	1	2.55	-73.952943	
1831 CMT		2012-06-21 23:38:42+00:00	2012-06-21 23:52:54+00:00	1	2.40	-73.984838	
1832 CMT		2012-07-10 12:08:36+00:00	2012-07-10 12:11:25+00:00	1	0.50	-74.004139	

```
In 1408 1 new_trips = new_trips.drop(['trip_distance'], axis=1)
2 new_trips
```

Out 1408

	vendor_id	pickup_datetime	dropoff_datetime	passenger_count	pickup_longitude	pickup_latitude	
0	CMT	2012-11-14 22:58:30+00:00	2012-11-14 23:20:28+00:00	1	-74.011095	40.708404	
1	VT	2012-09-07 12:22:00+00:00	2012-09-07 12:47:00+00:00	1	-73.963615	40.808487	
2	VT	2012-10-17 00:25:00+00:00	2012-10-17 00:45:00+00:00	4	-73.987095	40.748255	
4	CMT	2012-07-24 07:32:18+00:00	2012-07-24 07:45:48+00:00	1	-74.016593	40.709927	
5	VT	2012-05-24 00:47:00+00:00	2012-05-24 00:53:00+00:00	1	-74.000945	40.737432	
...	
1825	VT	2012-05-31 01:55:00+00:00	2012-05-31 02:07:00+00:00	3	-73.970242	40.762157	
1826	CMT	2012-03-05 10:37:12+00:00	2012-03-05 10:51:37+00:00	1	-73.979540	40.771234	
1827	VT	2012-03-05 10:37:12+00:00	2012-03-05 10:51:37+00:00	1	-73.979540	40.771234	

```
In 1409 1 #Discarding time stamp
2 new_trips['pickup_datetime'] = pd.to_datetime(new_trips['pickup_datetime'])
3 new_trips['pickup_datetime'] = new_trips['pickup_datetime'].apply(lambda x: x.timestamp()).astype(float)
4 new_trips['dropoff_datetime'] = pd.to_datetime(new_trips['dropoff_datetime'])
5 new_trips['dropoff_datetime'] = new_trips['dropoff_datetime'].apply(lambda x: x.timestamp()).astype(float)
```

```
In 1410 1 #new_trips['dropoff_datetime'] = pd.to_datetime(new_trips['dropoff_datetime']).dt.date
```

```
In 1411 1 new_trips = pd.get_dummies(new_trips, columns=['vendor_id'])
```

```
In 1412 1 #new_trips['pickup_datetime'] = pd.to_datetime(new_trips['pickup_datetime'])
2 #new_trips['pickup_datetime'] = new_trips['pickup_datetime'].astype(int)
```


In 1413 1 new_trips

Out 1413

	pickup_datetime	dropoff_datetime	passenger_count	pickup_longitude	pickup_latitude	dropoff_longitude	dropoff_latitude
0	1.352934e+09	1.352935e+09	1	-74.011095	40.708404	-73.984400	40.708404
1	1.347021e+09	1.347022e+09	1	-73.963615	40.808487	-74.005605	40.808487
2	1.350434e+09	1.350435e+09	4	-73.987095	40.748255	-74.006200	40.748255
4	1.343115e+09	1.343116e+09	1	-74.016593	40.709927	-73.972623	40.709927
5	1.337820e+09	1.337821e+09	1	-74.000945	40.737432	-73.984243	40.737432
...
1825	1.338429e+09	1.338430e+09	3	-73.970242	40.762157	-74.000485	40.762157
1826	1.330944e+09	1.330945e+09	1	-73.979540	40.771234	-73.971538	40.771234
1827	1.330944e+09	1.330945e+09	1	-73.979540	40.771234	-73.971538	40.771234

In 1414 1 new_trips.describe()

Out 1414

	pickup_datetime	dropoff_datetime	passenger_count	pickup_longitude	pickup_latitude	dropoff_longitude	dropoff_latitude
count	1.047000e+03	1.047000e+03	1047.000000	1047.000000	1047.000000	1047.000000	1047.000000
mean	1.341067e+09	1.341068e+09	1.729812	-73.940562	40.751926	-73.978336	40.751926
std	9.044042e+06	9.044051e+06	1.334823	1.636223	0.024091	0.022756	0.024091
min	1.325387e+09	1.325388e+09	1.000000	-74.017884	40.641342	-74.017936	40.641342
25%	1.333459e+09	1.333460e+09	1.000000	-73.993161	40.737179	-73.991677	40.737179
50%	1.340829e+09	1.340830e+09	1.000000	-73.983017	40.752625	-73.981428	40.752625
75%	1.348764e+09	1.348765e+09	2.000000	-73.971293	40.766622	-73.968218	40.766622
max	1.356998e+09	1.356998e+09	6.000000	-7.583332	40.847300	-73.702002	40.847300

```
In _ 1  
      2 new_trips.shape
```

```
Out: 1437 (1647, 13)
```

Before cleaning data we had 1833 rows 19 columns, but after cleaning we are left with 1647 rows and 13 columns. So total 186 rows have been removed.

(g) Split the quality-controlled data randomly into training (75%), validation (15%) and test (15%) sets. Show the descriptors of all these sets as in (d) of this question.

```
In 1416 1 X = new_trips.copy()
2 y = new_trips['total_amount']
```

```
In 1417 1 X_train, X_rem, y_train, y_rem = train_test_split(X,y, train_size=0.7)
```

```
In 1418 1 # Now since we want the valid and test size to be equal (15% each of overall data).
2 # we have to define valid_size=0.5 (that is 50% of remaining data)
3 test_size = 0.5
4 X_valid, X_test, y_valid, y_test = train_test_split(X_rem,y_rem, test_size=0.5)
5
6 print(X_train.shape), print(y_train.shape)
7 print(X_valid.shape), print(y_valid.shape)
8 print(X_test.shape), print(y_test.shape)
```

```
✓ (1152, 13)
  (1152,)
  (247, 13)
  (247,)
  (248, 13)
  (248,)
```

```
Out 1418 (None, None)
```

```
In 1419 1 X_train
```

In 1443 1 X_train.describe()

Out 1443 8 rows x 13 columns

	pickup_datetime	dropoff_datetime	passenger_count	pickup_longitude	pickup_latitude	dropoff_longitude	dropoff_latitude
count	1.152000e+03	1.152000e+03	1152.000000	1152.000000	1152.000000	1152.000000	1152.000000
mean	1.341043e+09	1.341043e+09	1.710938	-73.980761	40.751551	-73.978351	40.751551
std	9.077560e+06	9.077575e+06	1.317114	0.022346	0.023537	0.022801	0.023537
min	1.325387e+09	1.325388e+09	1.000000	-74.017884	40.641342	-74.017936	40.641342
25%	1.333350e+09	1.333350e+09	1.000000	-73.994070	40.737451	-73.991761	40.737451
50%	1.340698e+09	1.340699e+09	1.000000	-73.982891	40.752387	-73.981254	40.752387
75%	1.348924e+09	1.348925e+09	2.000000	-73.970906	40.766531	-73.967800	40.766531
max	1.356998e+09	1.356998e+09	6.000000	-73.776705	40.839140	-73.702002	40.839140

In 1444 1 y_train.describe()

Out 1444

```

count    1152.000000
mean      8.858854
std       3.995139
min       2.500000
25%       5.700000
50%       7.700000
75%      11.000000
max      22.300000
Name: total_amount, dtype: float64

```

```
In 1445 1 X_valid.describe()
```

Out 1445

	pickup_datetime	dropoff_datetime	passenger_count	pickup_longitude	pickup_latitude	dropoff_longitude	dropoff_latitude
count	2.470000e+02	2.470000e+02	247.000000	247.000000	247.000000	247.000000	247.000000
mean	1.340772e+09	1.340773e+09	1.842105	-73.980509	40.754850	-73.978192	40.754850
std	8.946191e+06	8.946197e+06	1.441039	0.017916	0.025613	0.022955	0.025613
min	1.325424e+09	1.325424e+09	1.000000	-74.017242	40.676129	-74.017012	40.676129
25%	1.333767e+09	1.333767e+09	1.000000	-73.991485	40.737580	-73.990928	40.737580
50%	1.340286e+09	1.340286e+09	1.000000	-73.983356	40.756167	-73.981683	40.756167
75%	1.348430e+09	1.348431e+09	2.000000	-73.970472	40.770292	-73.969877	40.770292
max	1.356685e+09	1.356686e+09	6.000000	-73.868980	40.847300	-73.785772	40.847300

```
In 1446 1 y_valid.describe()
```

Out 1446

count	247.000000
mean	8.821457
std	3.873601
min	2.500000
25%	6.000000
50%	8.000000
75%	11.150000
max	20.900000

Name: total_amount, dtype: float64

In 1447 1 X_test.describe()

Out 1447

	pickup_datetime	dropoff_datetime	passenger_count	pickup_longitude	pickup_latitude	dropoff_longitude	dropoff
count	2.480000e+02	2.480000e+02	248.000000	248.000000	248.000000	248.000000	
mean	1.341477e+09	1.341478e+09	1.705645	-73.714045	40.750759	-73.978413	
std	9.006710e+06	9.006701e+06	1.306341	4.216356	0.024944	0.022437	
min	1.325709e+09	1.325710e+09	1.000000	-74.017342	40.670012	-74.016106	
25%	1.333595e+09	1.333596e+09	1.000000	-73.993002	40.734846	-73.991897	
50%	1.341816e+09	1.341816e+09	1.000000	-73.983620	40.753335	-73.981730	
75%	1.348373e+09	1.348374e+09	2.000000	-73.973623	40.764276	-73.969567	
max	1.356872e+09	1.356872e+09	6.000000	-7.583332	40.840615	-73.862370	

In 1448 1 y_test.describe

Out 1448

```
<bound method NDFrame.describe of 84      5.7
398      6.5
519      9.0
781      5.3
990      5.0
...
840      6.1
1773     6.5
538      9.7
332      6.1
784      6.1
Name: total_amount, Length: 248, dtype: float64>
```

(h) Save the three data sets to CSV files and verify they were created and show some of their headers.

name: total_amount, length: 270, dtype: float64

```
In 1420 1 X_train.to_csv('mytraindata.csv', index=False)
```

```
In 1421 1 from IPython.display import FileLink
2 FileLink('mytraindata.csv')
```

Out 1421 mytraindata.csv

```
In 1422 1 X_train.head()
```

Out 1422 5 rows × 8 columns

	pickup_datetime	dropoff_datetime	passenger_count	pickup_longitude	pickup_latitude	dropoff_longitude	dropoff_latitude
1083	1.341239e+09	1.341239e+09	1	-74.000480	40.727296	-74.015109	40.727296
965	1.341219e+09	1.341219e+09	1	-73.960295	40.817830	-73.964845	40.817830
791	1.345055e+09	1.345055e+09	1	-73.983412	40.758132	-73.983412	40.758132
394	1.327341e+09	1.327343e+09	1	-73.989730	40.744926	-73.919194	40.744926
837	1.330356e+09	1.330356e+09	1	-73.066006	40.745256	-73.078668	40.745256

```
In 1422 1 y_train.to_csv('mytraindata_y.csv', index=False)
2 from IPython.display import FileLink
3
4 FileLink('mytraindata_y.csv')
```

Out 1422 mytraindata_y.csv

```
In 1450 1 y_train.head()
```

Out 1450 1083 7.7
965 4.9
791 3.3
394 18.1
837 11.3
Name: total_amount, dtype: float64

```
In 1423 1 X_valid.to_csv('myvaliddata.csv', index=False)  
2 from IPython.display import FileLink  
3  
4 FileLink('myvaliddata.csv')
```

Out 1423 myvaliddata.csv

```
In _ 1 X_valid.head()
```

Out 1451

5 rows x 13 columns								CSV				dropoff
#	pickup_datetime	dropoff_datetime	passenger_count	pickup_longitude	pickup_latitude	dropoff_longitude						
322	1.347906e+09	1.347906e+09	6	-73.963795	40.808062	-73.958015						
509	1.341763e+09	1.341764e+09	2	-73.975935	40.765498	-73.982147						
787	1.345653e+09	1.345654e+09	1	-73.956349	40.771674	-73.972545						
1616	1.355126e+09	1.355126e+09	1	-74.003900	40.738187	-73.994686						
1663	1.334752e+09	1.334753e+09	1	-73.081502	40.732592	-73.047221						

```
In 1424 1 y_valid.to_csv('myvaliddata_y.csv', index=False)
2 from IPython.display import FileLink
3
4 FileLink('myvaliddata_y.csv')
```

Out 1424 myvaliddata_y.csv

```
In 1452 1 y_valid.head()
```

Out 1452 ✓

322	4.5
509	6.1
787	8.5
1616	5.0
1443	13.7

Name: total_amount, dtype: float64

```
In 1425 1 X_test.to_csv('mytestdata.csv', index=False)
2 from IPython.display import FileLink
3
4 FileLink('mytestdata.csv')
```

Out 1425 mytestdata.csv

In 1453 1 X_test.head()

Out 1453 5 rows x 13 columns

	pickup_datetime	dropoff_datetime	passenger_count	pickup_longitude	pickup_latitude	dropoff_longitude	dropoff_latitude
84	1.327087e+09	1.327088e+09	4	-73.965140	40.755128	-73.969797	40.755128
398	1.353549e+09	1.353550e+09	2	-73.999690	40.738677	-73.998082	40.738677
519	1.350695e+09	1.350696e+09	1	-74.001615	40.740797	-73.982072	40.740797
781	1.326004e+09	1.326005e+09	1	-73.979458	40.735492	-73.987908	40.735492
990	1.351656e+09	1.351656e+09	1	-73.003753	40.754727	-73.003753	40.754727

In 1426 1 y_test.to_csv('mytestdata_y.csv', index=False)
 2 from IPython.display import FileLink
 3
 4 FileLink('mytestdata_y.csv')

Out 1426 mytestdata_y.csv

In _ 1 y_test.head()

Out 1454 84 5.7
 398 6.5
 519 9.0
 781 5.3
 990 5.0
 Name: total_amount, dtype: float64

(i) Use linear programming as well as other predictive models (such as regression or custom models) to predict the fare rate. Think of this machine learning model as a fare-estimation tool that could be put in a phone app to suggest a likely fare to New York Taxi riders so that they are not surprised and so that they can protest if the charge is much higher than expected.

```
In 1427 1 # load the training, validation, and test data
2 train_data = pd.read_csv('mytraindata.csv')
3 train_targets = pd.read_csv('mytraindata_y.csv')
4 valid_data = pd.read_csv('myvaliddata.csv')
5 valid_targets = pd.read_csv('myvaliddata_y.csv')
6 test_data = pd.read_csv('mytestdata.csv')
7 test_targets = pd.read_csv('mytestdata_y.csv')
```

```
In 1428 1 # fit the linear regression model to the training data
2 model = LinearRegression()
3 model.fit(train_data, train_targets)
```

Out 1428 LinearRegression()

```
In 1429 1 # evaluate the model on the validation set
2 valid_preds = model.predict(valid_data)
3 valid_mse = mean_squared_error(valid_targets, valid_preds)
4 print('Validation MSE:', valid_mse)
```

Validation MSE: 2.35310945942012e-21

```
.....
```

```
In 1430 1 # predict on the test set and calculate the test MSE
2 test_preds = model.predict(test_data)
3 test_mse = mean_squared_error(test_targets, test_preds)
4 print('Test MSE:', test_mse)
```

Test MSE: 2.3909759632081948e-21

```
In 1431 1 print(test_preds)
```

```

  - -
  ✓  [ 4.5]
    [ 8. ]
    [ 5.7]
    [ 5.5]
    [11.7]
    [10. ]
    [ 9.3]
    [13.5]
    [ 8.5]
    [10.1]
    [  . ]
```

5 6 7

```
In 1435 1 test_preds_df = pd.DataFrame(test_preds, columns=['predicted_total_amount'])
2
3 # Save the DataFrame to a CSV file
4 test_preds_df.to_csv('mypredicteddata.csv', index=False)
5
6 # Display a link to download the CSV file
7 from IPython.display import FileLink
8 FileLink('mypredicteddata.csv')
```

Out 1435 mypredicteddata.csv

I have used linear regression to predict data and attached all the training, validating, testing and predicted data csv files as well as python code for the verification purposes.

As we can see that mean square error is very small, so it shows that predicted total amount is very accurate. We can also compare in predicted data file and actual data file that total amount is very near to the actual amount.

NOTE: Linear Programming is generally not used to predict total amount for taxi fare as we need to predict continuous variable here. LP is used to optimize linear objective not for prediction. Therefore, linear regression works better for this problem.